

1. Read the Auto Data using Pandas

```
import pandas as pd
df = pd.read_csv("Auto.csv")
```

```
print(df.head(5))
print(df.shape)
```

```

      mpg  cylinders  displacement  horsepower  weight  acceleration  year  \
0   18.0         8         307.0         130    3504         12.0    70.0
1   15.0         8         350.0         165    3693         11.5    70.0
2   18.0         8         318.0         150    3436         11.0    70.0
3   16.0         8         304.0         150    3433         12.0    70.0
4   17.0         8         302.0         140    3449          NaN    70.0

```

```

      origin  name
0         1  chevrolet chevelle malibu
1         1      buick skylark 320
2         1  plymouth satellite
3         1      amc rebel sst
4         1      ford torino
(392, 9)

```

2. Data Exploration

```
df[["mpg", "weight", "year"]].describe()
# mpg -> Avg: 23.45 and Range: 37.6
# weight -> Avg: 2977.58 and Range: 3527
# year -> Avg: 76.01 and Range: 12
```

	mpg	weight	year
count	392.000000	392.000000	390.000000
mean	23.445918	2977.584184	76.010256
std	7.805007	849.402560	3.668093
min	9.000000	1613.000000	70.000000
25%	17.000000	2225.250000	73.000000
50%	22.750000	2803.500000	76.000000
75%	29.000000	3614.750000	79.000000
max	46.600000	5140.000000	82.000000

Saved successfully!



```
print('Before the changes:\n', df.dtypes)

df.cylinders = df.cylinders.astype('category').cat.codes
df.origin = df.origin.astype('category')

print('\nAfter the changes:\n', df.dtypes)
```

Before the changes:

```
mpg          float64
cylinders     int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        int64
name          object
dtype: object
```

After the changes:

```
mpg          float64
cylinders     int8
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        category
name          object
dtype: object
```

4. Delete Na's

```
df = df.dropna()
print('New Dimensions after NA drops: ', df.shape)
```

New Dimensions after NA drops: (389, 9)

5. Modify columns

```
avg_mpg = df.mpg.mean()

df['mpg_high'] = [1 if mpg > avg_mpg else 0 for mpg in df.mpg]
```

Saved successfully!



```
inplace=True)
print(df.head(5))
```

	cylinders	displacement	horsepower	weight	acceleration	year	origin	\
0	4	307.0	130	3504	12.0	70.0	1	
1	4	350.0	165	3693	11.5	70.0	1	
2	4	318.0	150	3436	11.0	70.0	1	
3	4	304.0	150	3433	12.0	70.0	1	
6	4	454.0	220	4354	9.0	70.0	1	

	mpg_high
0	0
1	0
2	0
3	0
6	0

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/10min/boolean_indexing.html
This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py:4913: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/10min/boolean_indexing.html
errors=errors,


6. Data exploration with graphs

```
import seaborn as sns
```

```
# with this graph, I learned that there fewer vehicles with a mpg higher than the average
g1 = sns.catplot(x='mpg_high', kind='count', data=df)
```




```
# this graphs shows that there is a direct correlation between weight and horsepower,
# this graphs also shows that the more weight/more horsepower has the more likely it is
g2 = sns.relplot(x='horsepower', y='weight', hue='mpg_high', data=df)
```



```
# this graph shows that cars that don't have high mileage tend to be much heavier than
g3 = sns.boxplot(x='mpg_high', y='weight', data=df)
```

Saved successfully!





7. Train/Test Split

```
from sklearn.model_selection import train_test_split

X = df.loc[:, ['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'mpg']]
y = df.mpg_high

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print('train size:', X_train.shape)
print('test size:', X_test.shape)

train size: (311, 7)
test size: (78, 7)
```

8. Logistic Regression

```
from sklearn.linear_model import LogisticRegression

lrm = LogisticRegression(max_iter=500, solver='lbfgs')

lrm.fit(X_train, y_train)
pred = lrm.predict(X_test)

# print metrics
from sklearn.metrics import classification_report

print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.84	0.91	0.87	50
1	0.84	0.91	0.87	28
avg / total	0.84	0.91	0.87	78

Saved successfully!



		1.00	0.88	28
accuracy			0.90	78
macro avg	0.89	0.92	0.89	78
weighted avg	0.92	0.90	0.90	78

9. Decision Tree

```
from sklearn import tree

dtm = tree.DecisionTreeClassifier()
dtm.fit(X_train, y_train)
pred = dtm.predict(X_test)

from sklearn.metrics import classification_report
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.96	0.92	0.94	50
1	0.87	0.93	0.90	28
accuracy			0.92	78
macro avg	0.91	0.92	0.92	78
weighted avg	0.93	0.92	0.92	78

10. Neural Network

```
from sklearn import preprocessing
from sklearn.neural_network import MLPClassifier

scaler = preprocessing.StandardScaler().fit(X_train)

X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

# train, test, evaluate first model
nnm = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5, 2), max_iter=500, random_state=1)
nnm.fit(X_train_scaled, y_train)
pred = nnm.predict(X_test)

from sklearn.metrics import classification_report
print(classification_report(y_test, pred))
```

```
# train, test, evaluate second model
```

Saved successfully!

gd', hidden_layer_sizes=(3,), max_iter=1500, random_state=1)

```

in)
pred = nnm_two.predict(X_test)

from sklearn.metrics import classification_report
print(classification_report(y_test, pred))

# Both models are the same, I think since our dataset is so small not only are neural

/usr/local/lib/python3.7/dist-packages/sklearn/base.py:444: UserWarning: X has feature names, but {self.__class__.__name__} was fitted without
f"X has feature names, but {self.__class__.__name__} was fitted without"
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UserWarning: Precision-Recall score is not available for multi-class
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UserWarning: Precision-Recall score is not available for multi-class
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UserWarning: Precision-Recall score is not available for multi-class
_warn_prf(average, modifier, msg_start, len(result))

```

	precision	recall	f1-score	support
0	0.64	1.00	0.78	50
1	0.00	0.00	0.00	28
accuracy			0.64	78
macro avg	0.32	0.50	0.39	78
weighted avg	0.41	0.64	0.50	78

	precision	recall	f1-score	support
0	0.64	1.00	0.78	50
1	0.00	0.00	0.00	28
accuracy			0.64	78
macro avg	0.32	0.50	0.39	78
weighted avg	0.41	0.64	0.50	78

```

/usr/local/lib/python3.7/dist-packages/sklearn/base.py:444: UserWarning: X has feature names, but {self.__class__.__name__} was fitted without
f"X has feature names, but {self.__class__.__name__} was fitted without"
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UserWarning: Precision-Recall score is not available for multi-class
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UserWarning: Precision-Recall score is not available for multi-class
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UserWarning: Precision-Recall score is not available for multi-class
_warn_prf(average, modifier, msg_start, len(result))

```

11. Analysis

- The Decision Tree algorithm performed the best.
-

11. Analysis

- The Decision Tree algorithm performed the best.
-

Saved successfully!



1. DT
2. LR
3. Both NN

Recall

1. Both NN
2. DT
3. LR

Precision

1. LR
2. DT
3. Both NN

c. Decision Trees and Logistic Regression are comparable to each other and both performed better than either Neural Networks. I think the reason this is because since the dataset is small the Neural Networks just have a huge disadvantage and are probably overfitting or learning from more noise than the other algorithms.

d. Honestly, they're similar and each have their own pros and cons. On one hand R has very intuitive and strict rules as to how everything needs to be done. There's only one way to do it but it makes sense once you find out what it is. Python is very different in that way because there are so many ways to accomplish the same task. I think I prefer Python, simply because there is more documentation and online help for it than there is for R.

Accuracy

1. DT
2. LR
3. Both NN

Recall

1. Both NN
2. DT
3. LR

Precision

1. LR
2. DT
3. Both NN

c. Decision Trees and Logistic Regression are very comparable to each other and both performed so much better than either Neural Networks. I think the reason this is because since the dataset is so small the Neural Networks just have a huge disadvantage and are probably overfitting or learning from more noise than the other algorithms. d. Honestly, they're similar and each have their own pros and cons. On one hand R has very intuitive and strict rules as to how everything needs to be done. There's only one way to do it but it makes sense once you find out what it is. Python is very different in that way because there are so many ways to accomplish the same task. I think I prefer Python, simply because there is more documentation and online help for it than there is for R.

Saved successfully!





Colab paid products - Cancel contracts here



Saved successfully!

X